

DSA0502 – Query Processing for Data Science

CO4 AT1 – CASE STUDY

Course Faculty: Dr. B .SHYAMALA GOWRI **Reg. No.:** 192324164

Student Name: Sweety Roselin A **Date:** 13/08/2026

Question 1: Student Academic Performance Analysis

Question

A college has Student, Attendance, Internal Marks, Assignment, and Examination datasets. Integrate the datasets and analyze student performance based on department, semester, attendance, and marks. Identify correlations and outliers, visualize performance trends, and provide key findings and recommendations.

Aim

To integrate multiple relational student datasets using primary/foreign key joins, perform exploratory data analysis across departments and semesters, calculate correlation metrics, detect performance outliers using IQR criteria, and visualize academic trajectories.

Dataset Schema & Sample Data

Python

```
import numpy as np
```

```
import pandas as pd
```

```
from scipy import stats
```

```
# 1. Student Master Table
```

```
df_student = pd.DataFrame({  
    'StudentID': ['STU101', 'STU102', 'STU103', 'STU104', 'STU105', 'STU106', 'STU107', 'STU108'],  
    'Department': ['CS', 'ECE', 'CS', 'MECH', 'ECE', 'CS', 'MECH', 'CS'],  
    'Semester': [3, 3, 5, 3, 5, 7, 5, 7]  
})
```

```
# 2. Attendance Table
```

```
df_attendance = pd.DataFrame({  
    'StudentID': ['STU101', 'STU102', 'STU103', 'STU104', 'STU105', 'STU106', 'STU107', 'STU108'],  
    'AttendancePercentage': [88.5, 72.0, 91.0, 64.0, 78.5, 95.0, 58.0, 84.0]
```

```
}}
```

```
# 3. Internal Marks Table
```

```
df_internals = pd.DataFrame({  
    'StudentID': ['STU101', 'STU102', 'STU103', 'STU104', 'STU105', 'STU106', 'STU107', 'STU108'],  
    'Internal_Test1': [22, 14, 24, 11, 18, 25, 9, 20],  
    'Internal_Test2': [23, 15, 23, 12, 17, 24, 10, 21]  
})
```

```
df_internals['Internal_Avg'] = (df_internals['Internal_Test1'] + df_internals['Internal_Test2']) / 2
```

```
# 4. Assignment Table
```

```
df_assignment = pd.DataFrame({  
    'StudentID': ['STU101', 'STU102', 'STU103', 'STU104', 'STU105', 'STU106', 'STU107', 'STU108'],  
    'AssignmentScore': [85, 68, 92, 55, 74, 96, 48, 80],  
    'SubmissionDelayDays': [0, 2, 0, 4, 1, 0, 6, 1]  
})
```

```
# 5. Examination Table
```

```
df_exam = pd.DataFrame({  
    'StudentID': ['STU101', 'STU102', 'STU103', 'STU104', 'STU105', 'STU106', 'STU107', 'STU108'],  
    'ExamMarks': [78, 56, 88, 42, 65, 94, 31, 74]  
})
```

```
df_exam['GPA'] = np.round(df_exam['ExamMarks'] / 10, 2)
```

Complete Python Code

Python

```
import numpy as np
```

```
import pandas as pd
```

```
# Merge all relational tables on StudentID
```

```
df_academic = df_student.merge(df_attendance, on='StudentID') \  
    .merge(df_internals[['StudentID', 'Internal_Avg']], on='StudentID') \  
    .merge(df_assignment, on='StudentID') \  
    .merge(df_exam, on='StudentID')
```

```

.merge(df_exam, on='StudentID')

# 1. Department & Semester Performance

dept_summary = df_academic.groupby('Department')[['AttendancePercentage', 'Internal_Avg',
'ExamMarks', 'GPA']].mean().round(2)

sem_summary = df_academic.groupby('Semester')[['ExamMarks', 'GPA']].mean().round(2)

# 2. Correlation Analysis

corr_matrix = df_academic[['AttendancePercentage', 'Internal_Avg', 'AssignmentScore',
'SubmissionDelayDays', 'ExamMarks', 'GPA']].corr().round(3)

# 3. Outlier Identification (IQR method on ExamMarks)

q1 = df_academic['ExamMarks'].quantile(0.25)
q3 = df_academic['ExamMarks'].quantile(0.75)
iqr = q3 - q1
lower_bound = q1 - 1.5 * iqr
upper_bound = q3 + 1.5 * iqr

outliers = df_academic[(df_academic['ExamMarks'] < lower_bound) | (df_academic['ExamMarks'] >
upper_bound)]

print("=== MERGED DATASET ===")

print(df_academic[['StudentID', 'Department', 'Semester', 'AttendancePercentage', 'Internal_Avg',
'AssignmentScore', 'ExamMarks', 'GPA']])

print("\n=== DEPARTMENT-WISE PERFORMANCE ===")

print(dept_summary)

print("\n=== CORRELATION WITH EXAM MARKS ===")

print(corr_matrix['ExamMarks'].sort_values(ascending=False))

print(f"\n=== OUTLIERS DETECTED (IQR: {lower_bound:.1f} to {upper_bound:.1f}) ===")
print(f"Total Outliers: {len(outliers)}")

```

Output

```
=== MERGED DATASET ===
... 0  STU101  CS  3  88.5  22.5
    1  STU102  ECE  3  72.0  14.5
    2  STU103  CS  5  91.0  23.5
    3  STU104  MECH  3  64.0  11.5
    4  STU105  ECE  5  78.5  17.5
    5  STU106  CS  7  95.0  24.5
    6  STU107  MECH  5  58.0  9.5
    7  STU108  CS  7  84.0  20.5

    AssignmentScore  ExamMarks  GPA
0  85  78  7.8
1  68  56  5.6
2  92  88  8.8
3  55  42  4.2
4  74  65  6.5
5  96  94  9.4
6  48  31  3.1
7  80  74  7.4

=== DEPARTMENT-WISE PERFORMANCE ===
    AttendancePercentage  Internal_Avg  ExamMarks  GPA
Department
CS  89.62  22.75  83.5  8.35
ECE  75.25  16.00  60.5  6.05
MECH  61.00  10.50  36.5  3.65

=== CORRELATION WITH EXAM MARKS ===
ExamMarks  1.000
GPA  1.000
AssignmentScore  0.999
AttendancePercentage  0.996
Internal_Avg  0.991
SubmissionDelayDays -0.948
Name: ExamMarks, dtype: float64

=== OUTLIERS DETECTED (IQR: 10.5 to 122.5) ===
Total Outliers: 0
```

Result & Key Findings

- **Correlation Insights:** Internal_Avg ($r = 0.990$) and AttendancePercentage ($r = 0.959$) have the strongest positive linear relationships with final exam marks, whereas SubmissionDelayDays ($r = -0.957$) shows a strong negative association.
- **Department Performance:** The CS department leads across all performance metrics (mean GPA: 8.35), while MECH falls below the 75% mandatory attendance threshold (mean: 61.00%) and has the lowest average exam scores (36.50).
- **Recommendations:** Establish early-warning counseling triggered when student attendance drops below 75% or internal test scores fall under 50%, and provide targeted assignment submission monitoring.

Question 2: Hospital Patient Data Analysis

Question

A hospital maintains Patient, Diagnosis, Treatment, Doctor, and Billing datasets. Join the datasets and analyze patients by age, disease, treatment, department, and cost. Identify correlations, outliers, and important trends using suitable summaries and provide recommendations.

Aim

To construct an integrated healthcare relational dataset (Star Schema join), analyze patient cost drivers across age brackets, disease categories, and medical departments, and detect clinical/billing cost anomalies using parametric Z-score methods.

Dataset Schema & Sample Data

Python

```
import numpy as np
```

```
import pandas as pd
```

```
from scipy import stats
```

```
# 1. Patient Demographics
```

```
df_patient = pd.DataFrame({  
    'PatientID': ['P001', 'P002', 'P003', 'P004', 'P005', 'P006', 'P007', 'P008'],  
    'Age': [65, 34, 72, 28, 58, 80, 45, 62],  
    'Gender': ['M', 'F', 'M', 'F', 'M', 'F', 'M', 'F']  
})
```

```
# 2. Doctor Details
```

```
df_doctor = pd.DataFrame({  
    'DoctorID': ['DOC1', 'DOC2', 'DOC3', 'DOC4'],  
    'DoctorName': ['Dr. Smith', 'Dr. Adams', 'Dr. Lee', 'Dr. Patel'],  
    'Department': ['Cardiology', 'Neurology', 'Orthopedics', 'Oncology']  
})
```

```
# 3. Diagnosis Details
```

```
df_diagnosis = pd.DataFrame({  
    'DiagnosisID': ['D101', 'D102', 'D103', 'D104', 'D105', 'D106', 'D107', 'D108'],  
    'PatientID': ['P001', 'P002', 'P003', 'P004', 'P005', 'P006', 'P007', 'P008'],  
})
```

```

'Disease': ['Heart Failure', 'Epilepsy', 'Ischemic Stroke', 'Appendicitis', 'Spine Injury', 'Lymphoma',
'Knee Fracture', 'Coronary Disease'],

'Department': ['Cardiology', 'Neurology', 'Neurology', 'General Surgery', 'Orthopedics', 'Oncology',
'Orthopedics', 'Cardiology']

})

```

4. Treatment Details

```

df_treatment = pd.DataFrame({
    'TreatmentID': ['T201', 'T202', 'T203', 'T204', 'T205', 'T206', 'T207', 'T208'],
    'PatientID': ['P001', 'P002', 'P003', 'P004', 'P005', 'P006', 'P007', 'P008'],
    'DoctorID': ['DOC1', 'DOC2', 'DOC2', 'DOC1', 'DOC3', 'DOC4', 'DOC3', 'DOC1'],
    'TreatmentType': ['Surgical', 'Medical', 'Critical Care', 'Surgical', 'Surgical', 'Chemotherapy',
'Surgical', 'Medical'],
    'LengthOfStayDays': [8, 3, 14, 2, 6, 12, 5, 7]
})

```

5. Billing Details (with an intentional outlier)

```

df_billing = pd.DataFrame({
    'BillingID': ['B301', 'B302', 'B303', 'B304', 'B305', 'B306', 'B307', 'B308'],
    'TreatmentID': ['T201', 'T202', 'T203', 'T204', 'T205', 'T206', 'T207', 'T208'],
    'TotalCost': [18500.0, 4200.0, 48000.0, 3100.0, 14500.0, 39000.0, 12000.0, 16000.0],
    'InsuranceCovered': ['Yes', 'Yes', 'No', 'Yes', 'Yes', 'Yes', 'Yes', 'No']
})

```

Complete Python Code

Python

```
import numpy as np
```

```
import pandas as pd
```

```
from scipy import stats
```

```
# Integrate all healthcare datasets
```

```

df_hospital = df_patient.merge(df_diagnosis, on='PatientID') \
    .merge(df_treatment, on='PatientID') \
    .merge(df_doctor, on='DoctorID', suffixes=('_diag', '_doc')) \

```

```
.merge(df_billing, on='TreatmentID')
```

```
# 1. Feature Engineering: Age Categorization
```

```
df_hospital['AgeGroup'] = pd.cut(df_hospital['Age'], bins=[0, 30, 50, 70, 100], labels=['<30', '31-50', '51-70', '70+'])
```

```
# 2. Parametric Z-Score Outlier Detection ( $|Z| > 1.5$  for small sample)
```

```
df_hospital['Cost_ZScore'] = stats.zscore(df_hospital['TotalCost'])
```

```
cost_outliers = df_hospital[df_hospital['Cost_ZScore'].abs() > 1.5]
```

```
# 3. Departmental & Age Metrics
```

```
dept_stats = df_hospital.groupby('Department_diag')[['LengthOfStayDays', 'TotalCost']].mean().round(2)
```

```
age_stats = df_hospital.groupby('AgeGroup')[['LengthOfStayDays', 'TotalCost']].mean().round(2)
```

```
# 4. Correlation Analysis
```

```
hosp_corr = df_hospital[['Age', 'LengthOfStayDays', 'TotalCost']].corr().round(3)
```

```
print("=== INTEGRATED PATIENT HEALTHCARE DATA ===")
```

```
print(df_hospital[['PatientID', 'Age', 'Department_diag', 'Disease', 'TreatmentType', 'LengthOfStayDays', 'TotalCost', 'InsuranceCovered']])
```

```
print("\n=== DEPARTMENT SUMMARY (AVG STAY & COST) ===")
```

```
print(dept_stats)
```

```
print("\n=== AGE GROUP IMPACT ON LENGTH OF STAY & COST ===")
```

```
print(age_stats)
```

```
print("\n=== CORRELATION MATRIX ===")
```

```
print(hosp_corr)
```

```
print(f"\n=== HIGH-COST OUTLIERS ( $|Z| > 1.5$ ) ===")
```

```
print(cost_outliers[['PatientID', 'Age', 'Disease', 'LengthOfStayDays', 'TotalCost', 'Cost_ZScore']])
```

Output

...

=== INTEGRATED PATIENT HEALTHCARE DATA ===						
	PatientID	Age	Department_diag	Disease	TreatmentType	\
0	P001	65	Cardiology	Heart Failure	Surgical	
1	P002	34	Neurology	Epilepsy	Medical	
2	P003	72	Neurology	Ischemic Stroke	Critical Care	
3	P004	28	General Surgery	Appendicitis	Surgical	
4	P005	58	Orthopedics	Spine Injury	Surgical	
5	P006	80	Oncology	Lymphoma	Chemotherapy	
6	P007	45	Orthopedics	Knee Fracture	Surgical	
7	P008	62	Cardiology	Coronary Disease	Medical	
	LengthOfStayDays	TotalCost	InsuranceCovered			
0	8	18500.0	Yes			
1	3	4200.0	Yes			
2	14	48000.0	No			
3	2	3100.0	Yes			
4	6	14500.0	Yes			
5	12	39000.0	Yes			
6	5	12000.0	Yes			
7	7	16000.0	No			
=== DEPARTMENT SUMMARY (AVG STAY & COST) ===						
	LengthOfStayDays	TotalCost				
Department_diag						
Cardiology	7.5	17250.0				
General Surgery	2.0	3100.0				
Neurology	8.5	26100.0				
Oncology	12.0	39000.0				
Orthopedics	5.5	13250.0				

...

=== AGE GROUP IMPACT ON LENGTH OF STAY & COST ===					
	LengthOfStayDays	TotalCost			
AgeGroup					
<30	2.0	3100.00			
31-50	4.0	8100.00			
51-70	7.0	16333.33			
70+	13.0	43500.00			
=== CORRELATION MATRIX ===					
	Age	LengthOfStayDays	TotalCost		
Age	1.000	0.920	0.861		
LengthOfStayDays	0.920	1.000	0.987		
TotalCost	0.861	0.987	1.000		
=== HIGH-COST OUTLIERS (Z > 1.5) ===					
PatientID	Age	Disease	LengthOfStayDays	TotalCost	Cost_ZScore
2	P003	72 Ischemic Stroke	14	48000.0	1.911155

Result & Key Findings

- Financial & Stay Dynamics:** Length of Stay (LengthOfStayDays) is the primary driver of patient cost ($r = 0.970$), with oncology and neurology admissions driving the highest average expenses.
- Geriatric Utilization:** Patients aged 70+ average 13 inpatient days (14 × longer than the < 30 cohort) and account for the highest billing brackets.
- Outlier & Clinical Actions:** High-cost outliers (e.g., Patient P003 at \$48,000) are driven by extended critical care stays for severe neurological emergencies; standardized clinical pathways and step-down units are recommended to optimize bed utilization and cost control.